

# CS101 Data Structures

Binary Search Trees

Textbook Ch 12

# Outline

- Sorted list ADT
- Binary search tree
  - Definition
  - Implementation
- AVL tree
- Red-black tree
- More.

# Outline

This topic covers binary search trees:

- Abstract Sorted Lists
- Background
- Definition and examples
- Implementation:
  - Front, back, insert, erase
  - Previous smaller and next larger objects
  - Finding the  $k^{\text{th}}$  object

# Sorted List ADT

Previously, we discussed Abstract Lists

- the objects are explicitly ordered by the programmer

We will now discuss the Abstract Sorted List:

- the objects are ordered by their values

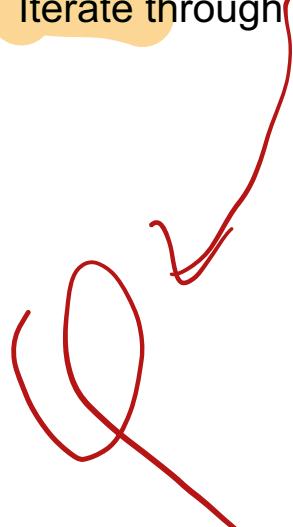
Certain operations no longer make sense:

- ~~push\_front~~ and ~~push\_back~~ are replaced by a generic **insert**

# Sorted List ADT

Queries that can be made about data in a Sorted List ADT include:

- Finding the smallest and largest entries
- Finding the  $k^{\text{th}}$  largest entry
- Find the next larger and previous smaller objects of a given object which may or may not be in the container
- Iterate through those objects that fall on an interval  $[a, b]$



# Implementation

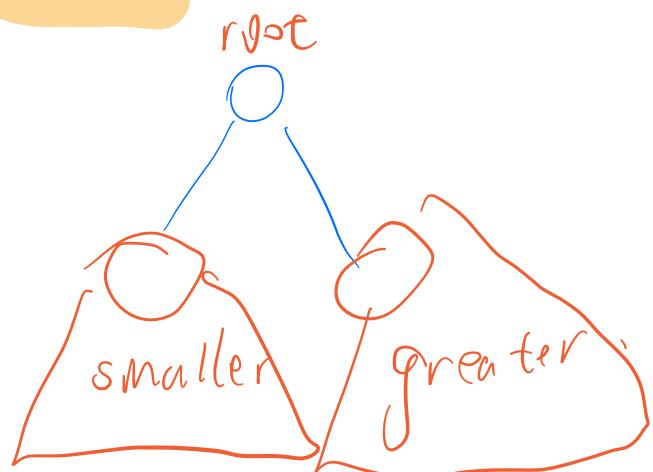
If we implement an Abstract Sorted List using an array or a linked list, we will have operations which are  **$O(n)$**

- As an insertion could occur anywhere in a linked list or array, we must either traverse or copy, on average,  **$O(n)$**  objects

# Binary Search Trees

In a binary search tree, we require that

- all objects in the left sub-tree to be less than the object stored in the root node
- all objects in the right sub-tree to be greater than the object in the root object
- the two sub-trees are themselves binary search trees



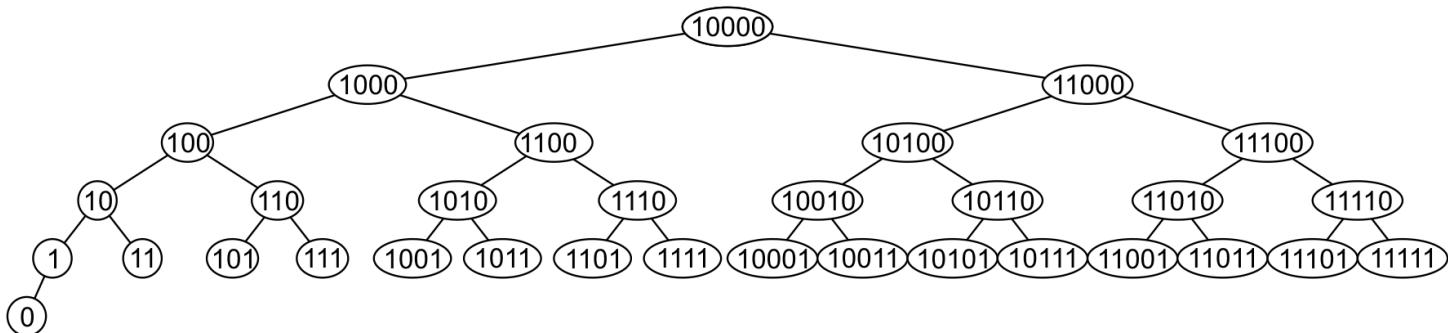
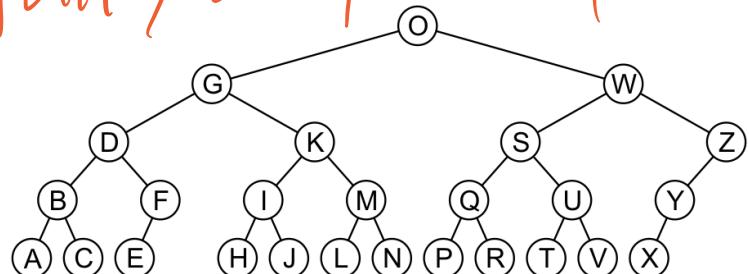
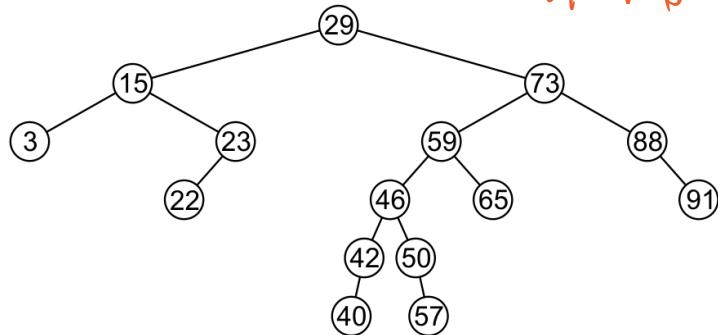
# Definition

Thus, we define a non-empty binary search tree as a binary tree with the following properties:

- The left sub-tree (if any) is a binary search tree and all elements are less than the root element, and
- The right sub-tree (if any) is a binary search tree and all elements are greater than the root element

# Examples

不~~用~~ full/completed



# Search

To search an object: examine the root node and if we have not found what we are looking for:

- If the object is less than what is stored in the root node, continue searching in the left sub-tree
- Otherwise, continue searching the right sub-tree

Time complexity:

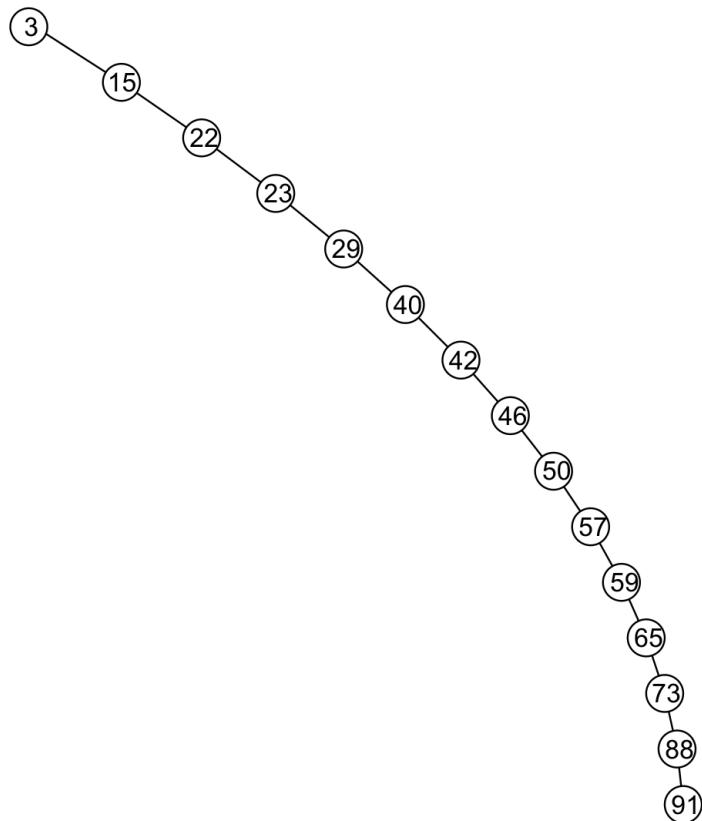
- $\mathbf{O}(h)$

# Worst case

Unfortunately, it is possible to construct *degenerate* binary search trees

壓迫 .

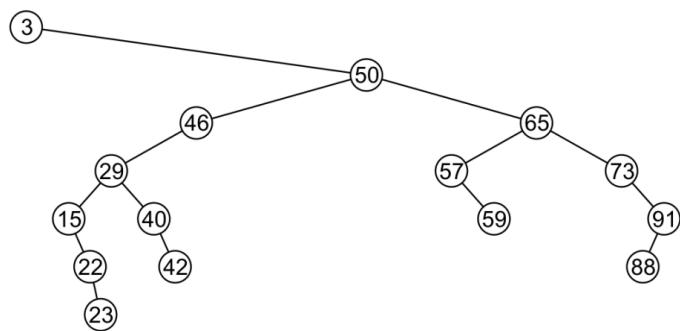
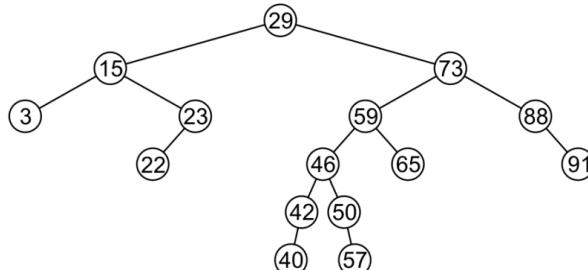
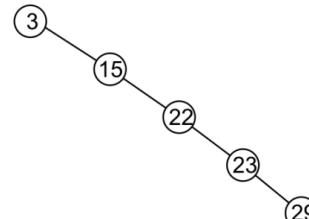
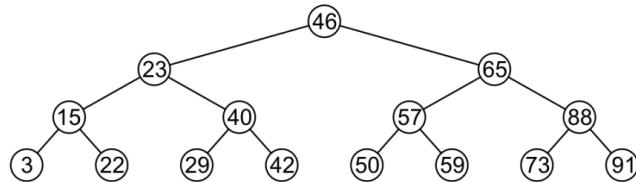
- This is equivalent to a linked list, i.e.,  $\mathbf{O}(n)$



# Examples

Order diff

All these binary search trees store the same data



# Duplicate Elements

4  
5  
3  
1

We will assume that in any binary tree, we are not storing duplicate elements unless otherwise stated

- In reality, it is seldom the case where duplicate elements in a container must be stored as separate entities

You can always consider duplicate elements with modifications to the algorithms we will cover

# Implementation

We will look at an implementation of a binary search tree in the same spirit as we did with our **Single\_list** class

- We will have a `Binary_search_nodes` class
- A `Binary_search_tree` class will store a pointer to the root

We will use templates, however, we will require that the class overrides the comparison operators

# Implementation

Any class which uses this binary-search-tree class must therefore implement:

```
bool operator<=( Type const &, Type const & );
bool operator< ( Type const &, Type const & );
bool operator==( Type const &, Type const & );
```

That is, we are allowed to compare two instances of this class

- Examples: **int** and **double**

# Implementation

```
#include "Binary_node.h"
```

```
template <typename Type>
class Binary_search_node: public Binary_node<Type> {
```

```
    using Binary_node<Type>::element;
```

```
    using Binary_node<Type>::left_tree;
```

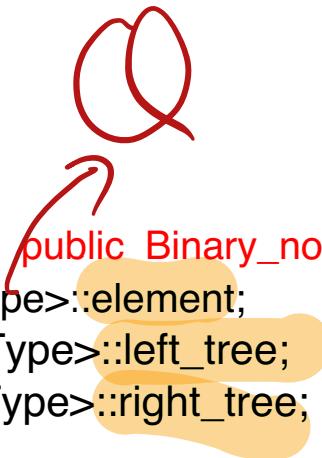
```
    using Binary_node<Type>::right_tree;
```

```
public:
```

```
    Binary_search_node( Type const & );
```

```
    Binary_search_node *left() const;
```

```
    Binary_search_node *right() const;
```



# Implementation

*node find ?*

```
Type front() const;  
Type back() const;  
bool find( Type const & ) const;  
  
bool insert( Type const & );  
bool erase( Type const &, Binary_search_node *& );  
};
```



# Constructor

The constructor simply calls the constructor of the base class

- Recall that it sets both `left_tree` and `right_tree` to `nullptr`
- It assumes that this is a new leaf node

```
template <typename Type>
Binary_search_node<Type>::Binary_search_node( Type const &obj ):
    Binary_node<Type>( obj ) {
    // Just calls the constructor of the base class
}
```

# Inherited Member Functions

The member functions

```
Type retrieve() const;  
bool is_leaf() const  
int size() const  
int height() const
```

are inherited from the base class Binary\_node

# left(), right()

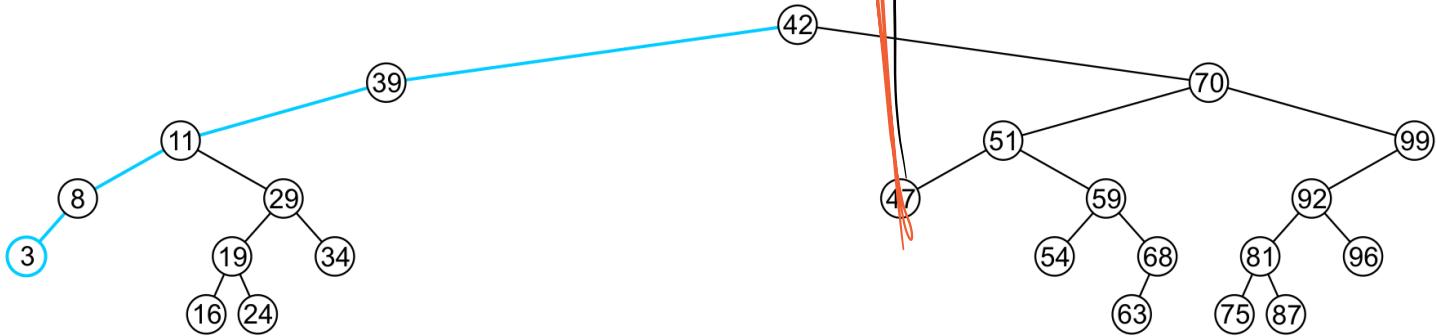
The base class returns a pointer to a `Binary_node`, we must recast them as `Binary_search_node`:

```
template <typename Type>
Binary_search_node<Type> *Binary_search_node<Type>::left() const {
    return reinterpret_cast<Binary_search_node *>(
        Binary_node<Type>::left() );
}

template <typename Type>
Binary_search_node<Type> *Binary_search_node<Type>::right() const {
    return reinterpret_cast<Binary_search_node *>(
        Binary_node<Type>::right() );
}
```

# Finding the Minimum Object

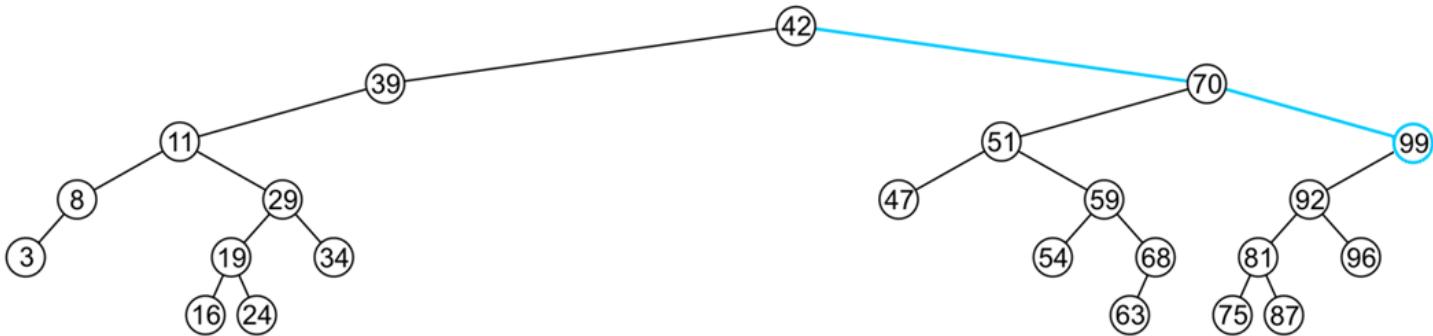
```
template <typename Type>
Type Binary_search_node<Type>::front() const {
    return ( left() == nullptr ) ? retrieve() : left()->front();
}
```



- The run time  $\mathbf{O}(h)$

# Finding the Maximum Object

```
template <typename Type>
Type Binary_search_node<Type>::back() const {
    return ( right() == nullptr ) ? retrieve() : right()->back();
}
```

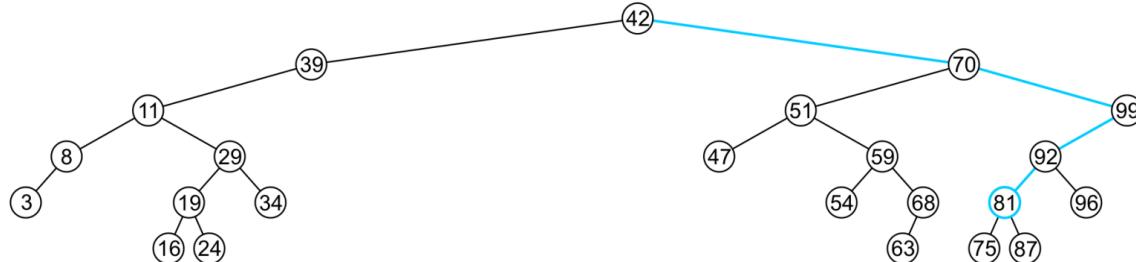


- The extreme values are not necessarily leaf nodes

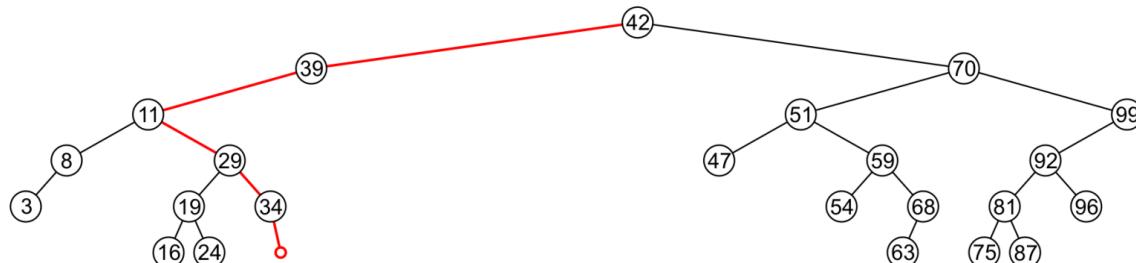
# Find

To determine membership, traverse the tree based on the linear relationship:

- If a node containing the value is found, e.g., 81, return true



- If an empty node is reached, e.g., 36, the object is not in the tree:



# Find

The implementation is similar to **front** and **back**:

```
template <typename Type>
bool Binary_search_node<Type>::find( Type const &obj ) const {
    if ( retrieve() == obj ) {
        return true;
    }

    if( obj < retrieve() )
        return left()==nullptr? false : left()->find( obj );
    else
        return right()==nullptr? false : right()->find( obj );
}
```

- The run time is  $O(h)$

*Worse =  $O(n)$*



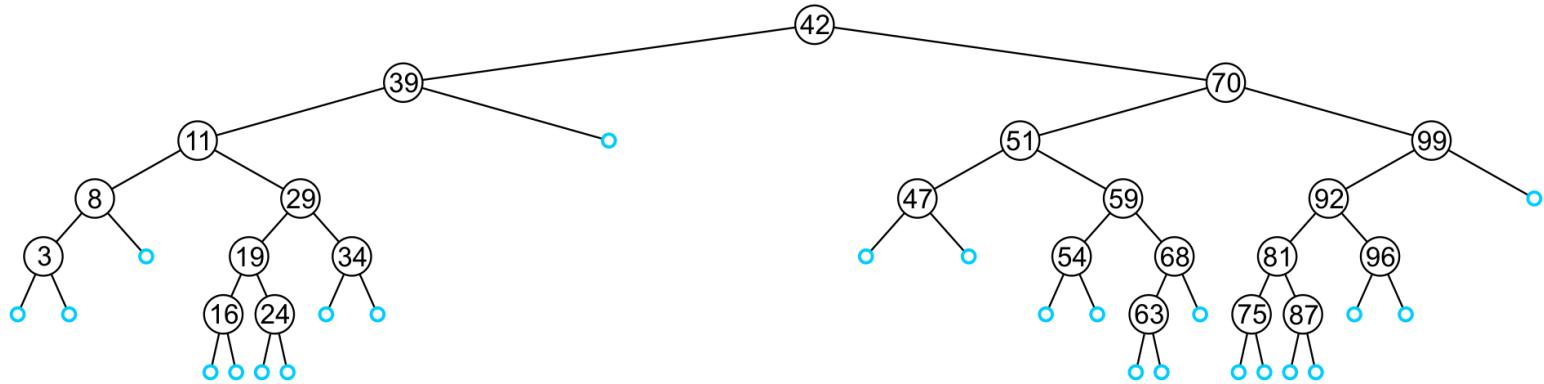
# Insert

Recall that a Sorted List is implicitly ordered

- It does not make sense to have member functions such as `push_front` and `push_back`
- Insertion will be performed by a single `insert` member function which places the object into the correct location

# Insert

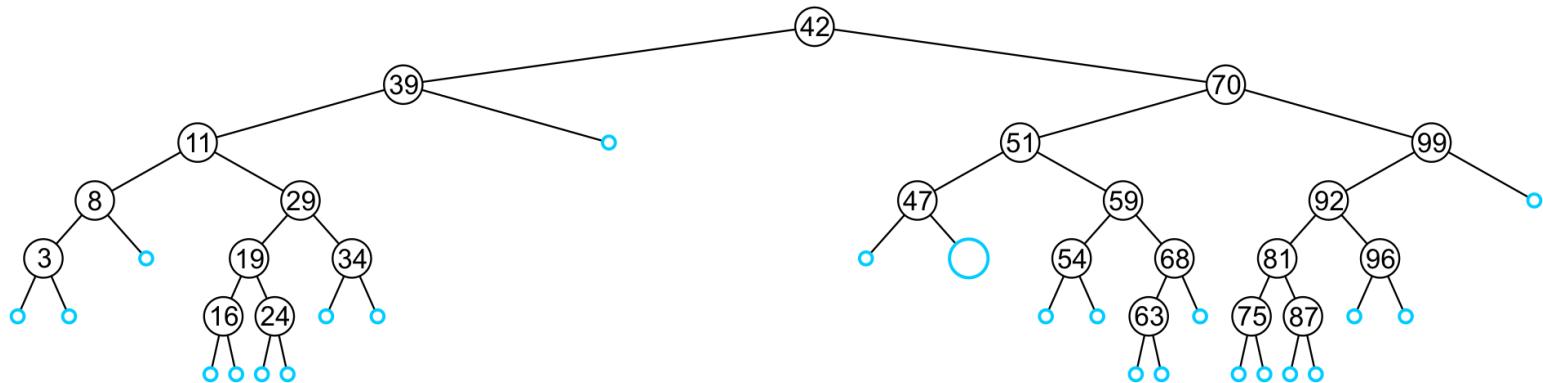
Any empty node is a possible location for an insertion



The values which may be inserted at any empty node depend on the surrounding nodes

# Insert

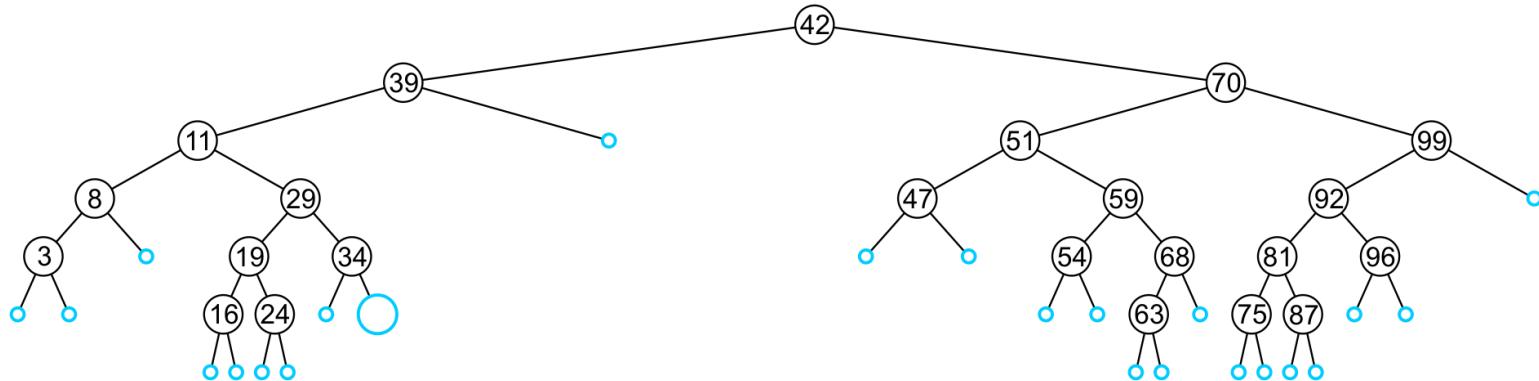
For example, this node may hold 48, 49, or 50



# Insert

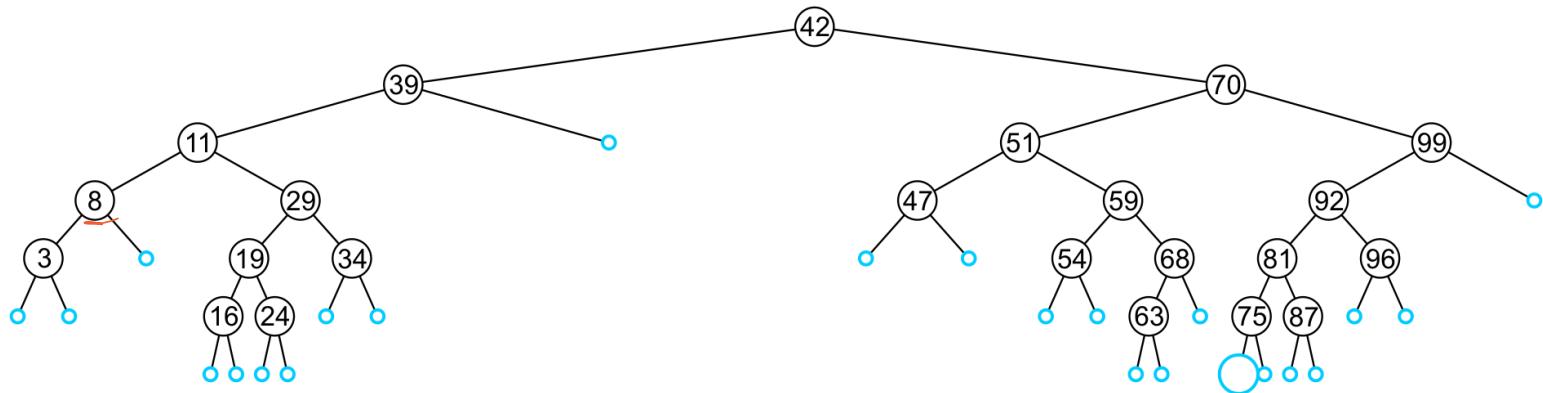
*limited range*

An insertion at this location must be 35, 36, 37, or 38



# Insert

This empty node may hold values from 71 to 74



# Insert

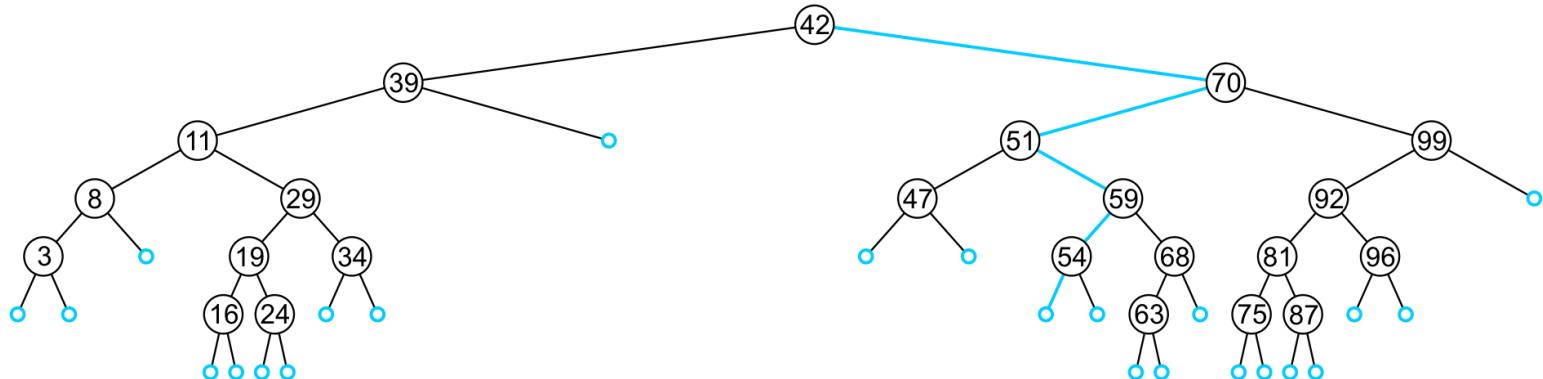
Like find, we will step through the tree

- If we find the object already in the tree, we will return
  - The object is already in the binary search tree (no duplicates)
- Otherwise, we will arrive at an empty node
- The object will be inserted into that location
- The run time is  $O(h)$

# Insert

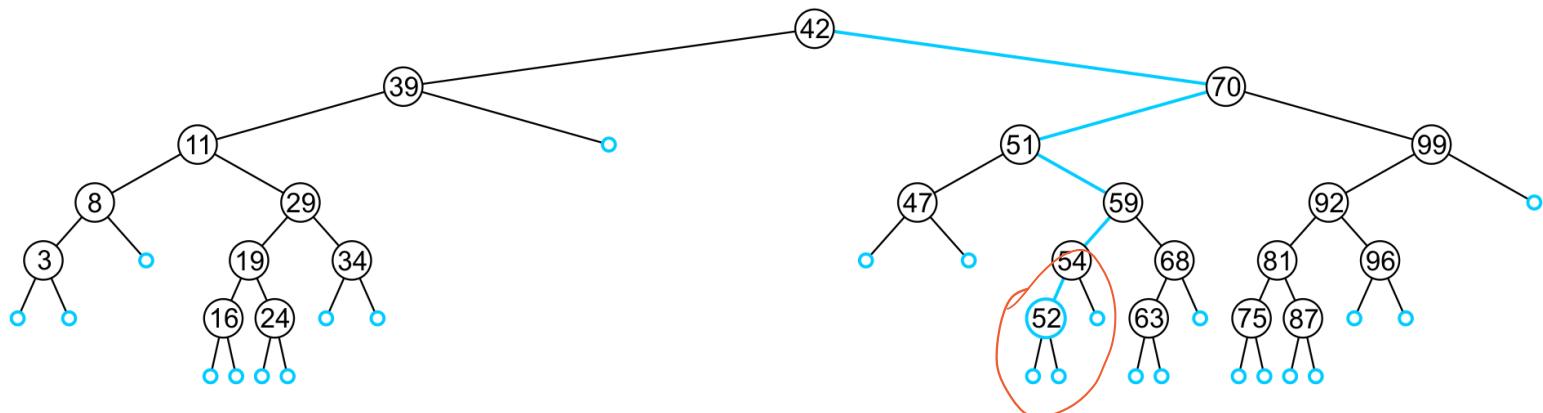
In inserting the value 52, we traverse the tree until we reach an empty node

- The left sub-tree of 54 is an empty node



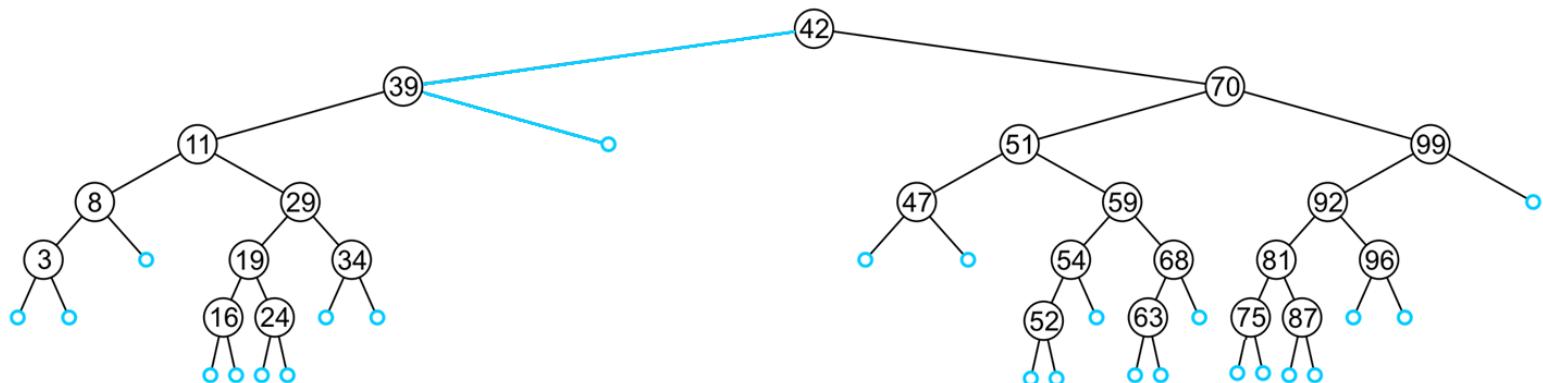
# Insert

A new leaf node is created and assigned to the member variable  
**left\_tree**



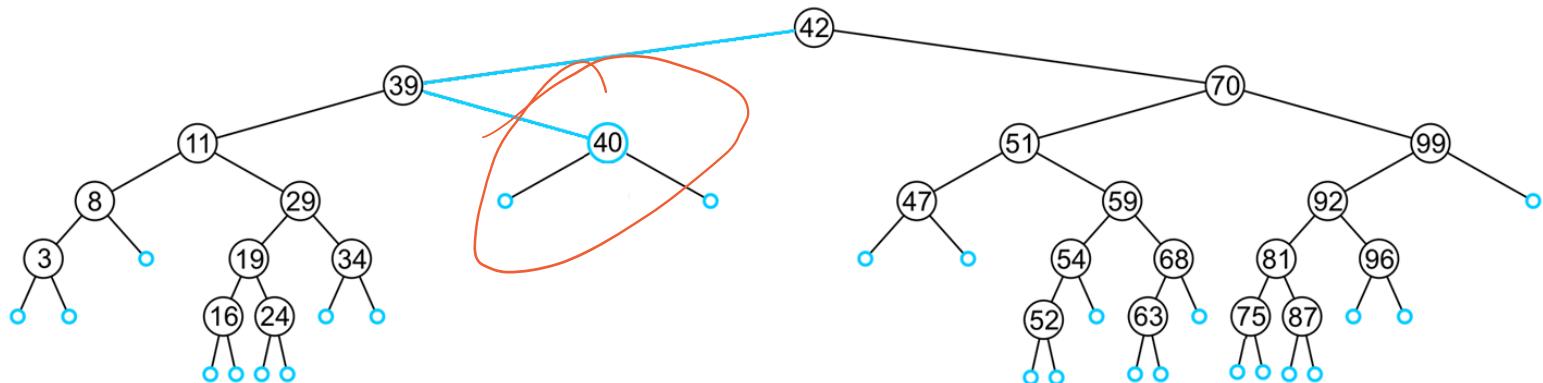
# Insert

In inserting 40, we determine the right sub-tree of 39 is an empty node

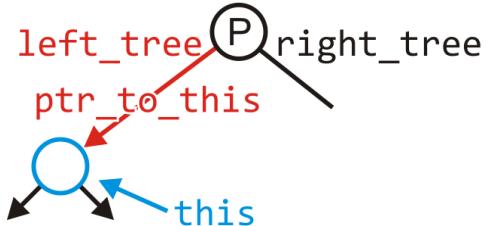


# Insert

A new leaf node storing 40 is created and assigned to the member variable `right_tree`



# Insert



```
template <typename Type>
bool Binary_search_node<Type>::insert( Type const &obj,
                                         Binary_search_node *ptr_to_this ) {
    if ( empty() ) {
        ptr_to_this = new Binary_search_node<Type>( obj );
        return true;
    } else if ( obj < retrieve() ) {
        return left()->insert( obj, left_tree );
    } else if ( obj > retrieve() ) {
        return right()->insert( obj, right_tree );
    } else {
        return false;
    }
}
```

将为子树  
的 right/  
left

# Insert

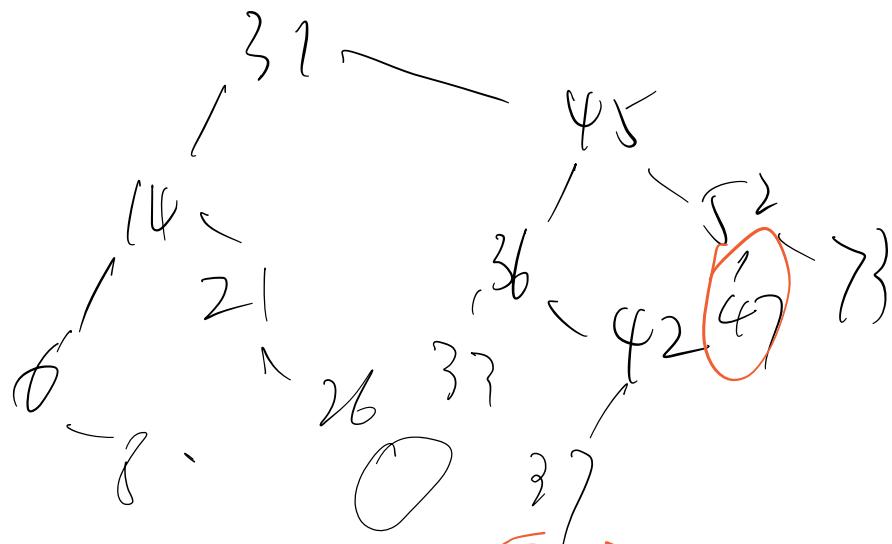
It is assumed that if neither of the conditions:

$\text{obj} < \text{retrieve}()$

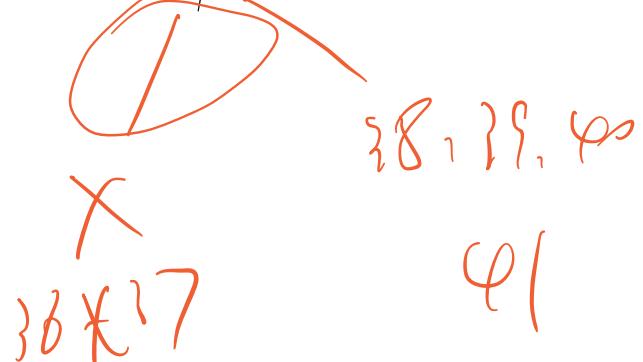
$\text{obj} > \text{retrieve}()$

then  $\text{obj} == \text{retrieve}()$  and therefore we do nothing

- The object is already in the binary search tree



# Insert



Blackboard example:

- In the given order, insert these objects into an initially empty binary search tree:

31 45 36 14 52 42 6 21 73 47 26 37 33 8

- What values could be placed:

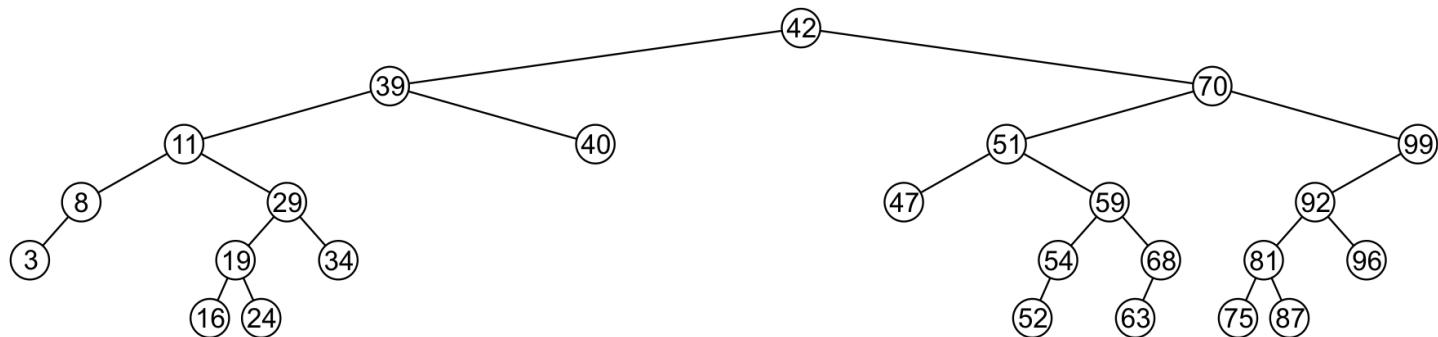
- To the left of 21? 15 - 20,
- To the right of 26? 27 - 30 .
- To the left of 47? 46 ,

- How would we determine if 40 is in this binary search tree?
- Which values could be inserted to increase the height of the tree?

# Erase

There are three possible scenarios:

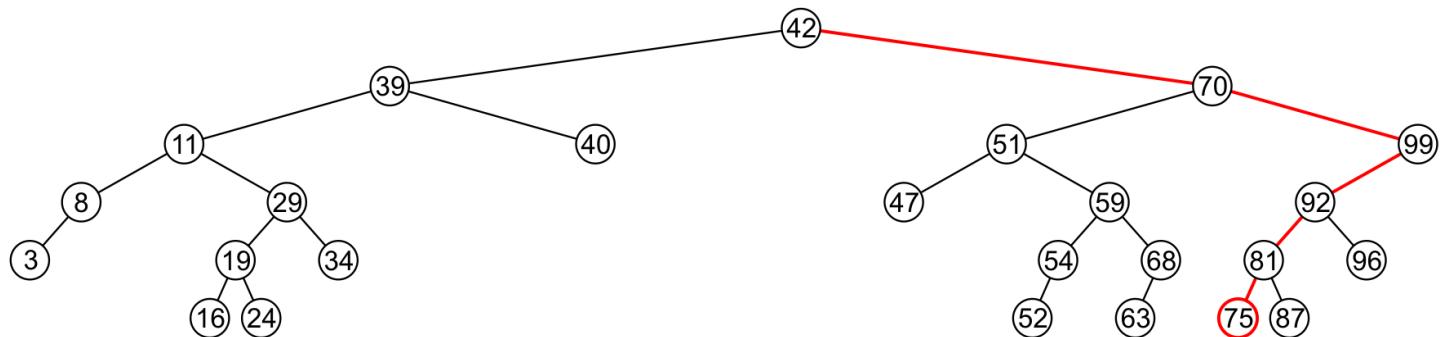
- The node is a leaf node,
- It has exactly one child, or
- It has two children (it is a full node)



# Erase

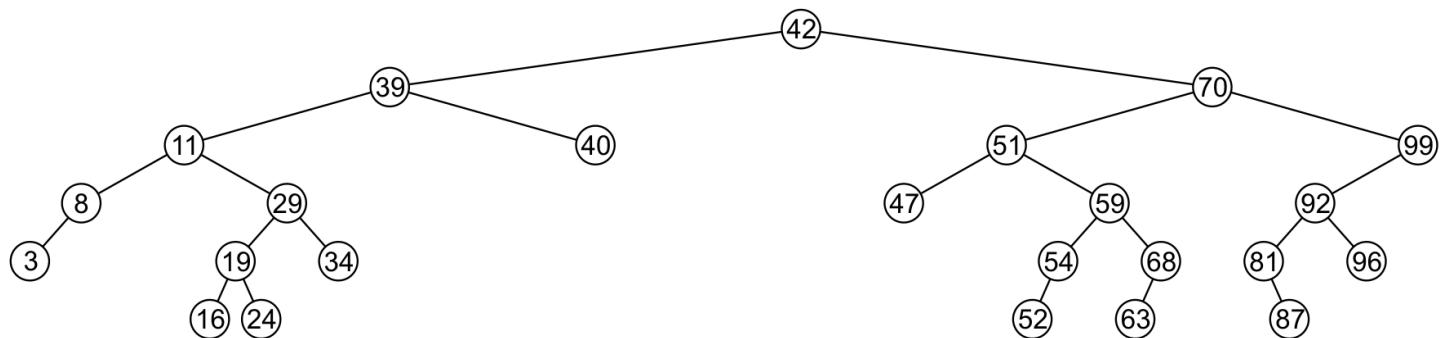
A leaf node simply must be removed and the appropriate member variable of the parent is set to **nullptr**

- Consider removing 75



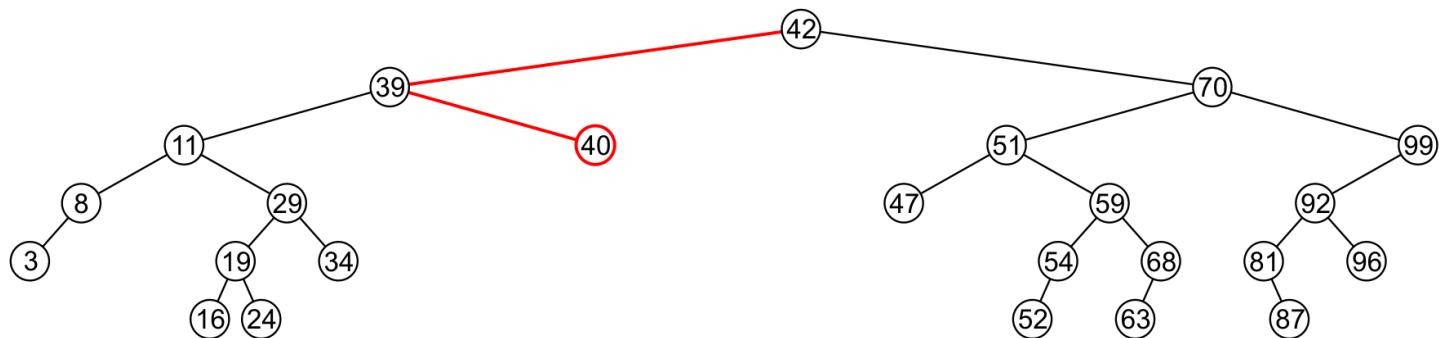
# Erase

The node is deleted and `left_tree` of 81 is set to `nullptr`



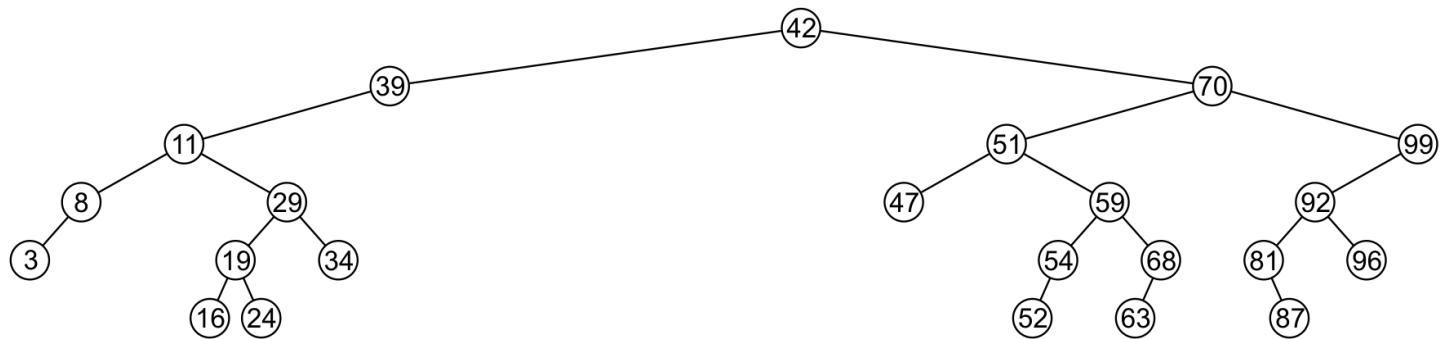
# Erase

Erasing the node containing 40 is similar



# Erase

The node is deleted and `right_tree` of 39 is set to `nullptr`

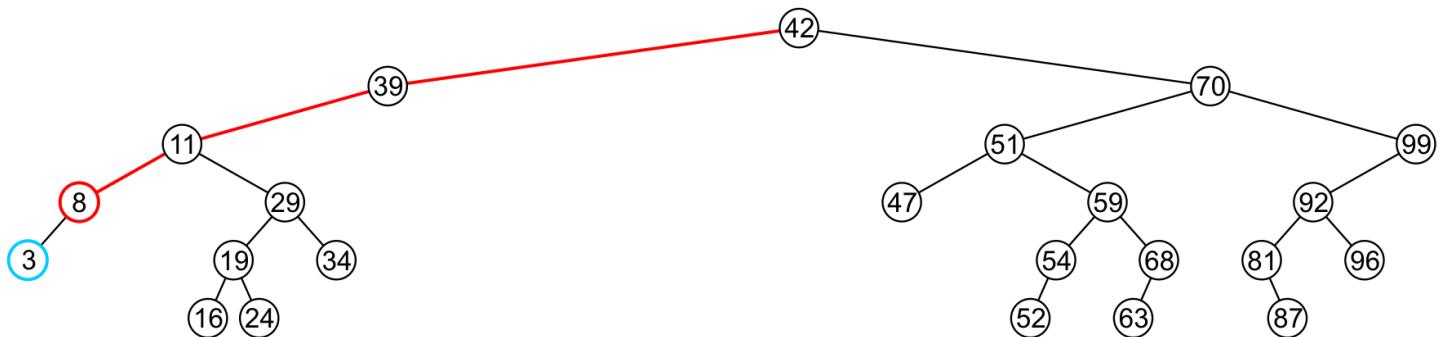


# Erase

If a node has only one child...

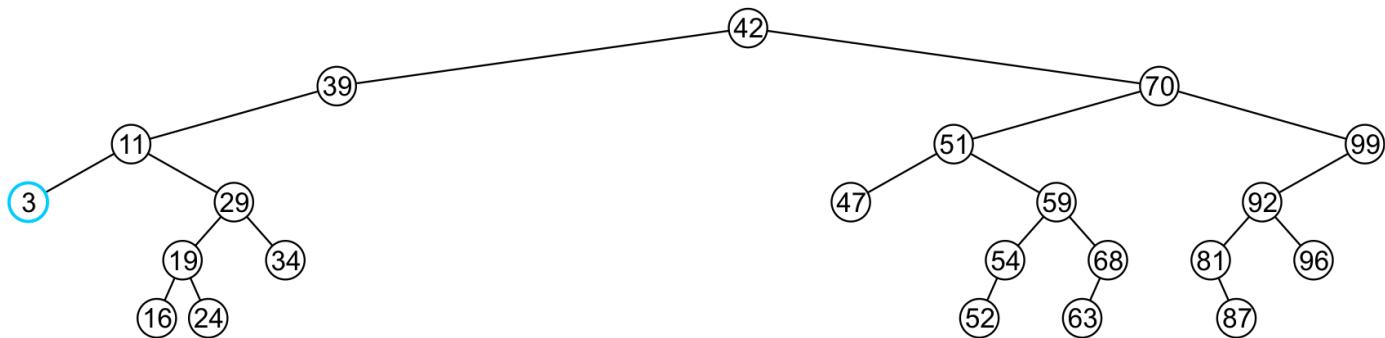
we can simply promote the sub-tree associated with the child

- Consider removing 8 which has one left child



# Erase

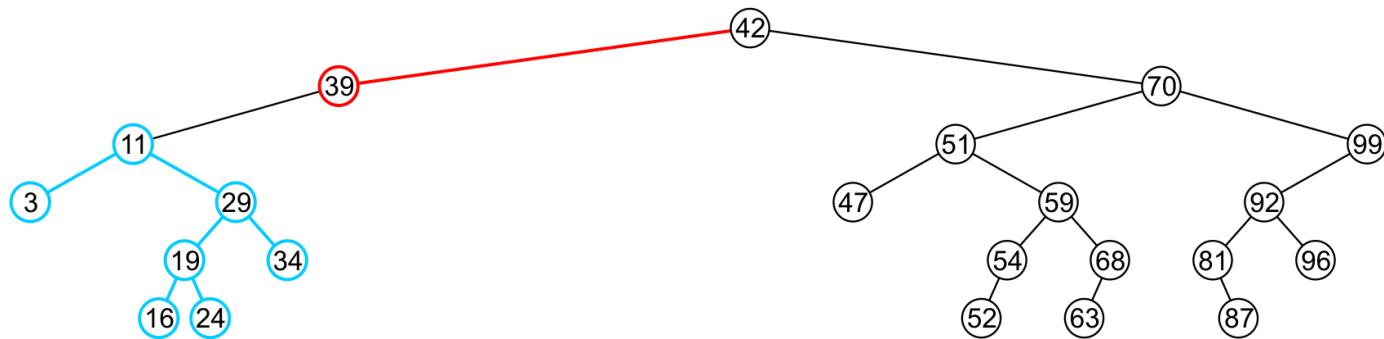
The node 8 is deleted and the `left_tree` of 11  
is updated to point to 3



# Erase

There is no difference in promoting a single node or a sub-tree

- To remove 39, it has a single child 11

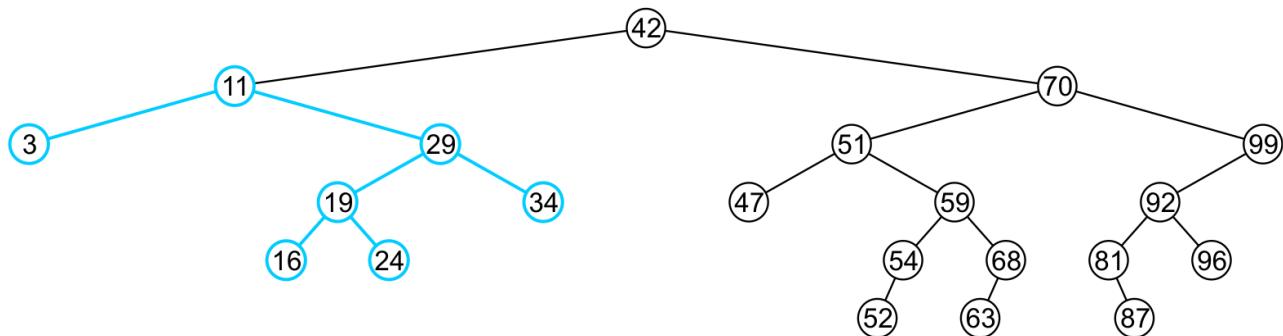


order mainta

# Erase

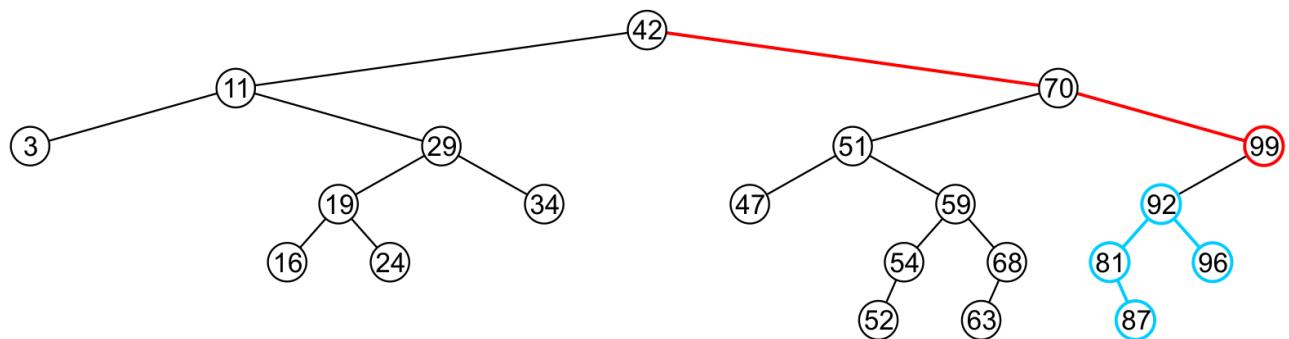
The node containing 39 is deleted and `left_node` of 42 is updated to point to 11

- Notice that order is still maintained



# Erase

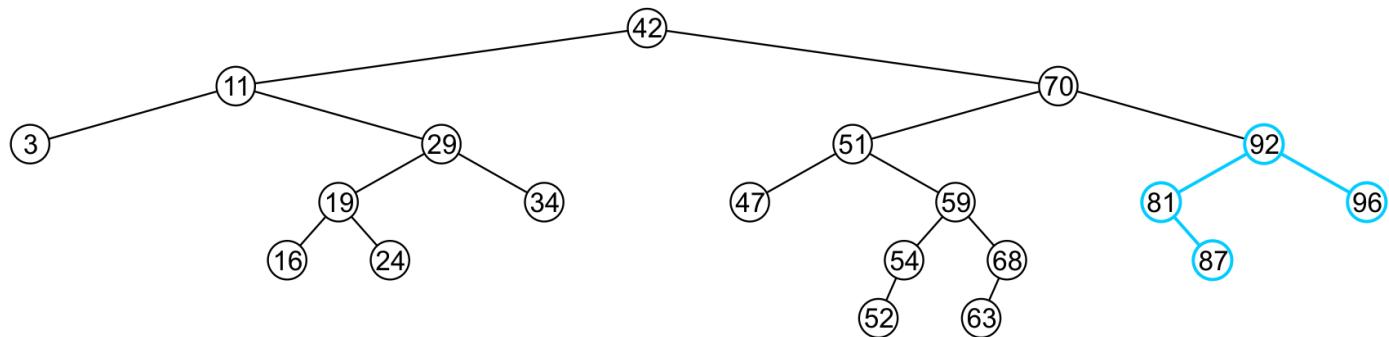
Consider erasing the node containing 99



# Erase

The node is deleted and the left sub-tree is promoted:

- The member variable `right_tree` of 70 is set to point to 92
- Again, the order of the tree is maintained



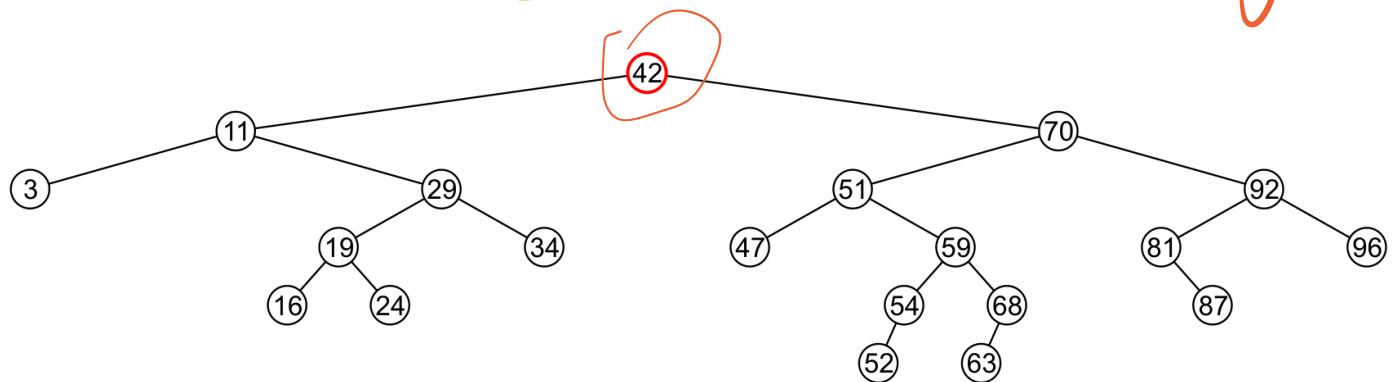
# Erase

Finally, we will consider the problem of erasing a full node, e.g., 42

We will perform two operations:

- Replace 42 with the minimum object in the right sub-tree
- Erase that object from the right sub-tree

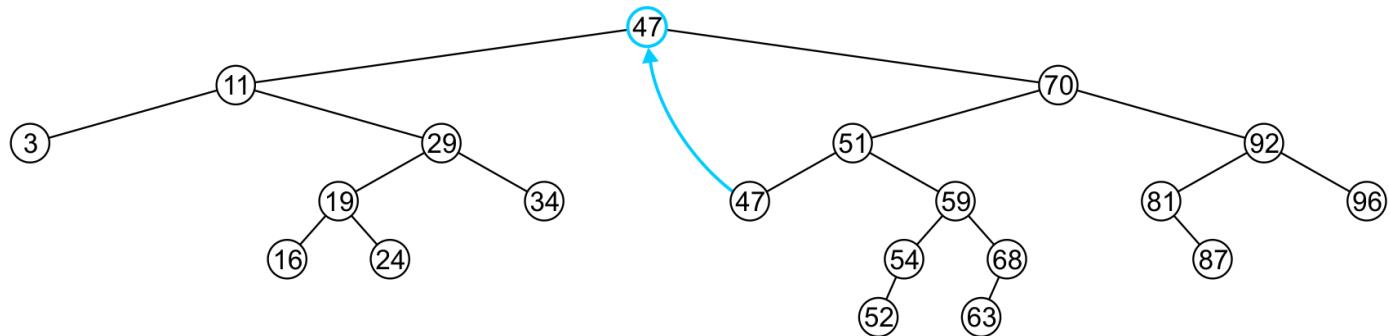
Right



# Erase

In this case, we replace 42 with 47

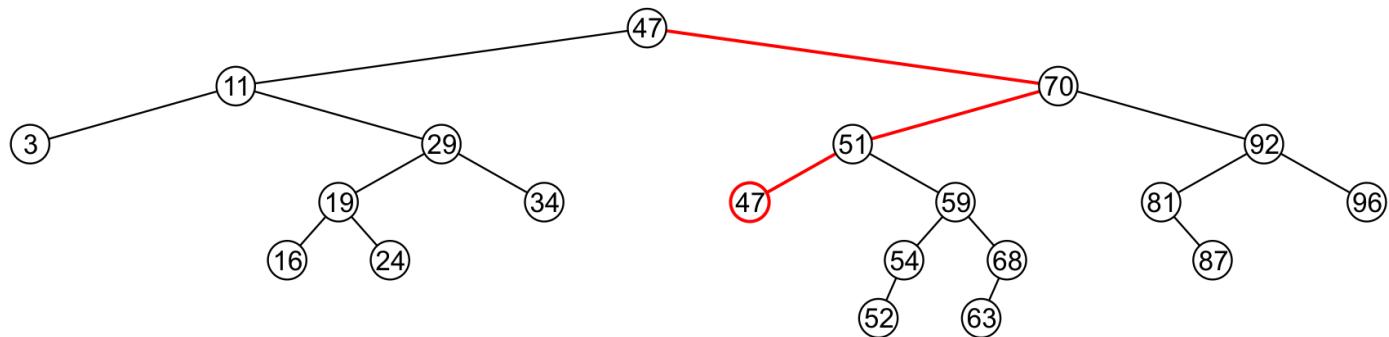
- We temporarily have two copies of 47 in the tree



# Erase

We now recursively erase 47 from the right sub-tree

- We note that 47 is a leaf node in the right sub-tree

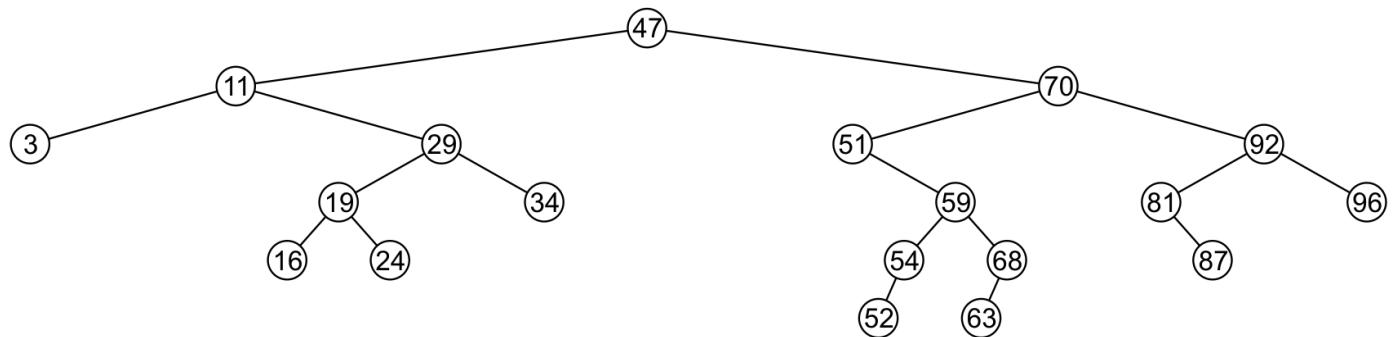


# Erase

Leaf nodes are simply removed and `left_tree` of 51 is set to `nullptr`

- Notice that the tree is still sorted:

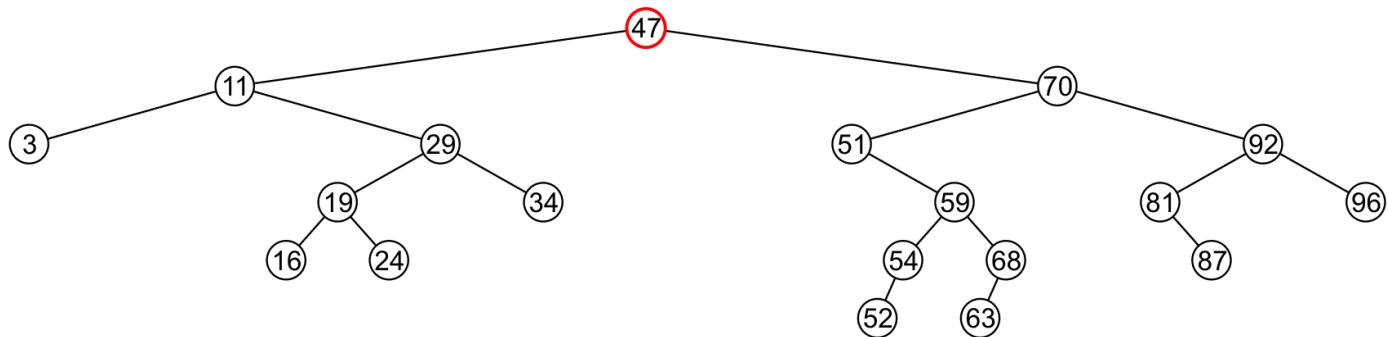
47 was the least object in the right sub-tree



# Erase

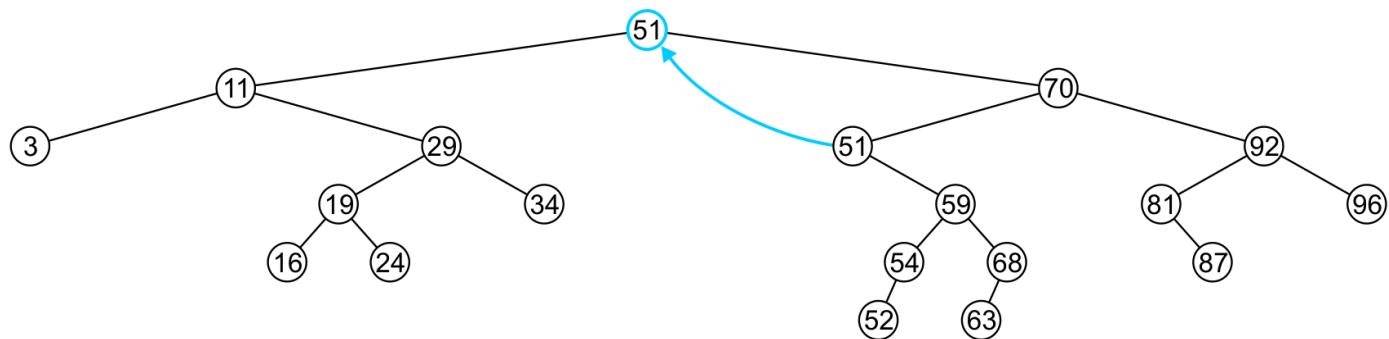
Suppose we want to erase the root 47 again:

- We must copy the minimum of the right sub-tree
- We could promote the maximum object in the left sub-tree and achieve similar results



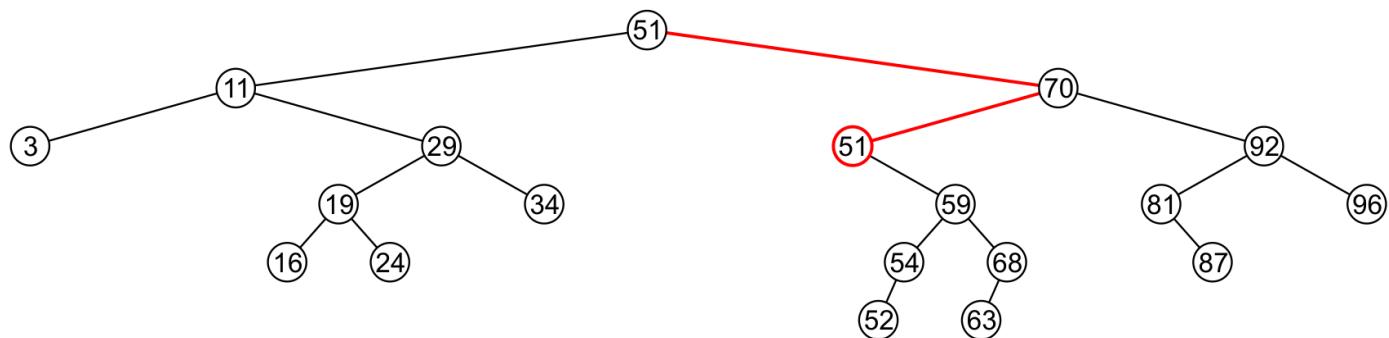
# Erase

We copy 51 from the right sub-tree



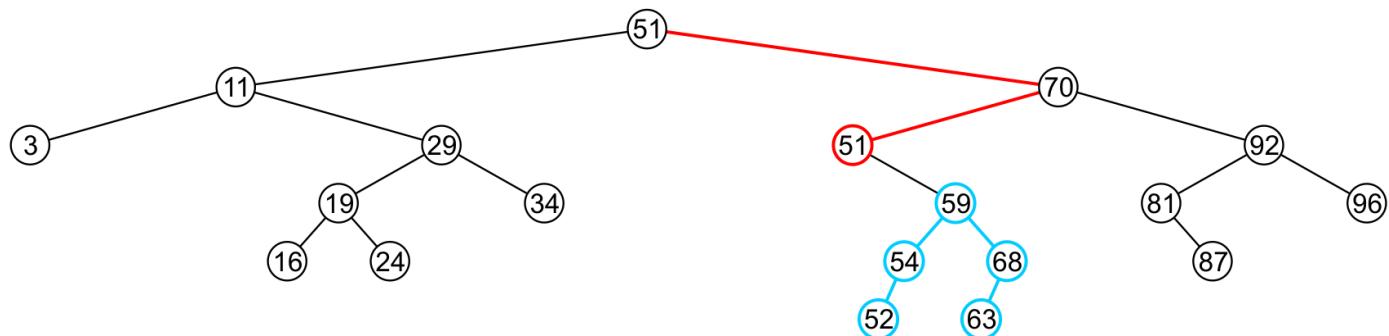
# Erase

We must proceed by delete 51 from the right sub-tree



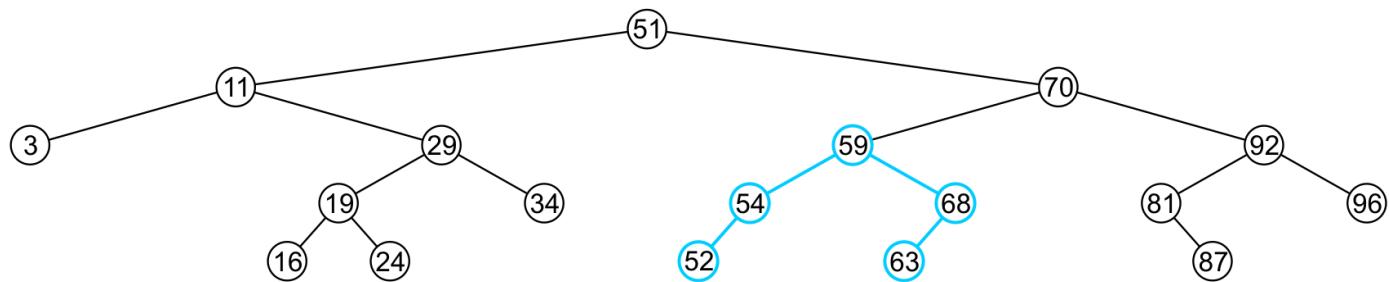
# Erase

In this case, the node storing 51 has just a single child



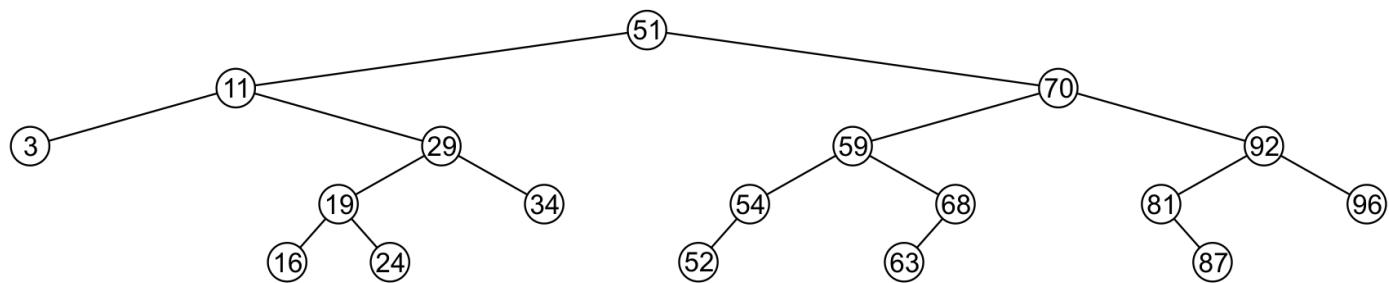
# Erase

We delete the node containing 51 and assign the member variable **left\_tree** of 70 to point to 59



# Erase

Note that after seven removals, the remaining tree is still correctly sorted



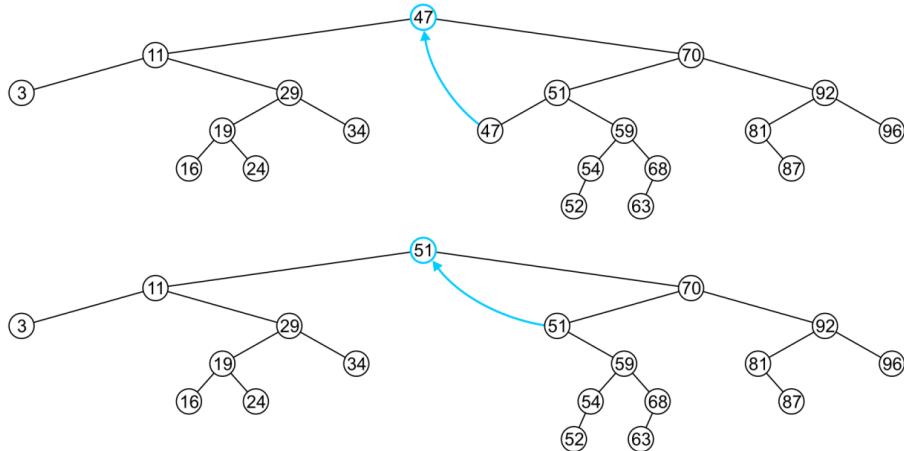
# Erase

In the two examples of removing a full node, we promoted:

- A node with no children
- A node with right child

What about a node with two children?

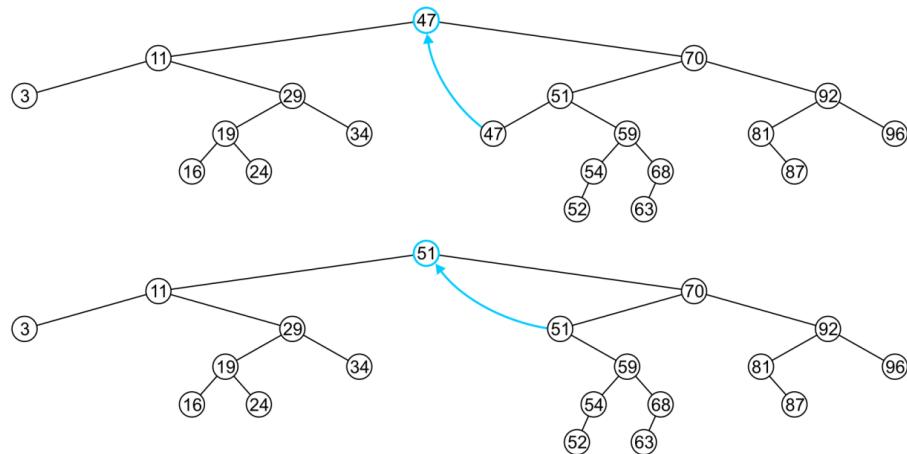
It is impossible for the node to have two children



# Erase

Recall that we promoted the minimum element in the right sub-tree

- If that node had a left sub-tree, that sub-tree would contain a smaller value



# Erase

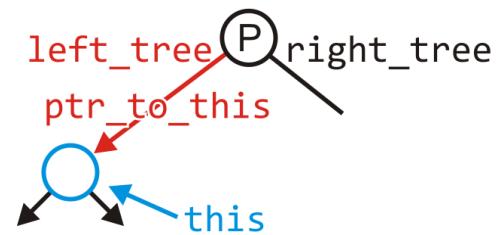
In order to properly remove a node, we will have to change the member variable pointing to the node

- To do this, we will pass that member variable by reference

Additionally: We will return 1 if the object is removed and 0 if the object was not found

# Erase

```
template <typename Type>
bool Binary_search_node<Type>::erase( Type const &obj, Binary_search_node *&ptr_to_this ) {
    if ( empty() ) {
        return false;
    } else if ( obj ==retrieve() ) {
        if ( is_leaf() ) { // leaf node
            ptr_to_this = nullptr;
            delete this;
        } else if ( !left()->empty() && !right()->empty() ) { // full node
            element = right()->front();
            right()->erase(retrieve(), right_tree );
        } else { // only one child
            ptr_to_this = ( !left()->empty() ) ? left() : right();
            delete this;
        }
        return true;
    } else if ( obj <retrieve() ) {
        return left()->erase( obj, left_tree );
    } else {
        return right()->erase( obj, right_tree );
    }
}
```



# Erase

Blackboard example:

- In the binary search tree generated previously:
  - Erase 47
  - Erase 21
  - Erase 45
  - Erase 31
  - Erase 36

$\mathcal{O}(h)$

$\mathcal{O}(n)$

# Binary Search Tree

We have defined binary search nodes

- Similar to the Single\_node in Project 1

We must now introduce a container which stores the root

- A Binary\_search\_tree class

Most operations will be simply passed to the root node

# Implementation

```
template <typename Type>
class Binary_search_tree {
private:
    Binary_search_node<Type> *root_node;
    Binary_search_node<Type> *root() const;
public:
    Binary_search_tree();
    ~Binary_search_tree();

    bool empty() const;
    int size() const;
    int height() const;
    Type front() const;
    Type back() const;
    int count( Type const &obj ) const;

    void clear();
    bool insert( Type const &obj );
    bool erase( Type const &obj );
};
```

# Constructor, Destructor, and Clear

```
template <typename Type>
Binary_search_tree<Type>::Binary_search_tree():
root_node( nullptr ) {
    // does nothing
}

template <typename Type>
Binary_search_tree<Type>::~Binary_search_tree() {
    clear();
}

template <typename Type>
void Binary_search_tree<Type>::clear() {
    root()->clear( root_node );
}
```

# Constructor, Destructor, and Clear

```
template <typename Type>
Binary_search_tree<Type> *Binary_search_tree<Type>::root() const {
    return tree_root;
}

template <typename Type>
bool Binary_search_tree<Type>::empty() const {
    return root()->empty();
}

template <typename Type>
int Binary_search_tree<Type>::size() const {
    return root()->size();
}
```

# Empty, Size, Height and Count

```
template <typename Type>
int Binary_search_tree<Type>::height() const {
    return root()->height();
}
```

```
template <typename Type>
bool Binary_search_tree<Type>::find( Type const &obj ) const {
    return root()->find( obj );
}
```

frost  
back

# Front and Back

// If root() is nullptr, 'front' will throw an underflow exception

template <typename Type>

Type Binary\_search\_tree<Type>::front() const {  
 return root()->front();

}

// If root() is nullptr, 'back' will throw an underflow exception

template <typename Type>

Type Binary\_search\_tree<Type>::back() const {  
 return root()->back();

}

# Insert and Erase

```
template <typename Type>
bool Binary_search_tree<Type>::insert( Type const &obj ) {
    return root()->insert( obj, root_node );
}

template <typename Type>
bool Binary_search_tree<Type>::erase( Type const &obj ) {
    return root()->erase( obj, root_node );
}
```

# Other Relation-based Operations

We will quickly consider two other relation-based queries that are very quick to calculate with an array of sorted objects:

- Finding the previous and next entries, and
- Finding the  $k^{\text{th}}$  entry

# Previous and Next Objects

Operations specific to linearly ordered data include:

- Find the next (larger) and previous (smaller) objects of a given object which may or may not be in the sorted list
- Find the  $k^{\text{th}}$  entry of the sorted list
- Iterate through those objects that fall on an interval  $[a, b]$

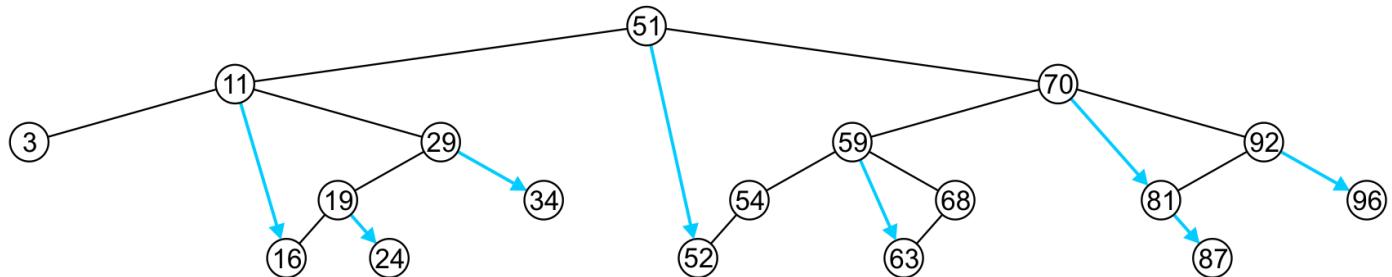
We will focus on finding the next (larger) object

- The others will follow

# Previous and Next Objects

To find the next object:

- If the node has a right sub-tree, the minimum object in that sub-tree is the next object

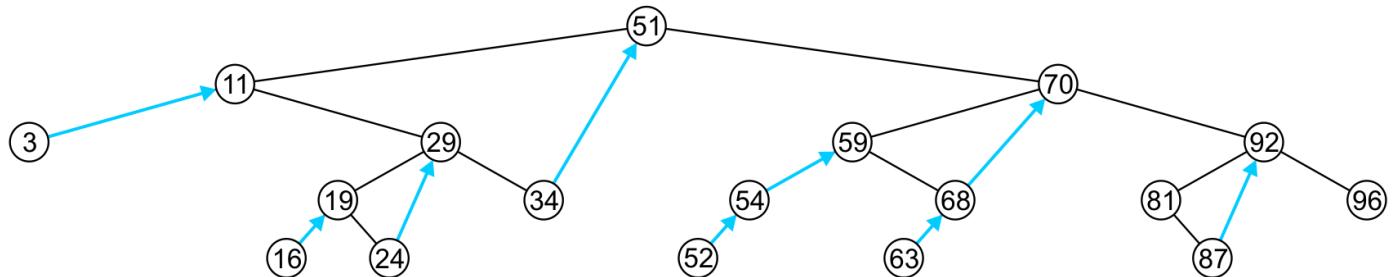


# Previous and Next Objects

B

If, however, there is no right sub-tree:

- It is the first larger object (if any) that exists in the path from the node to the root

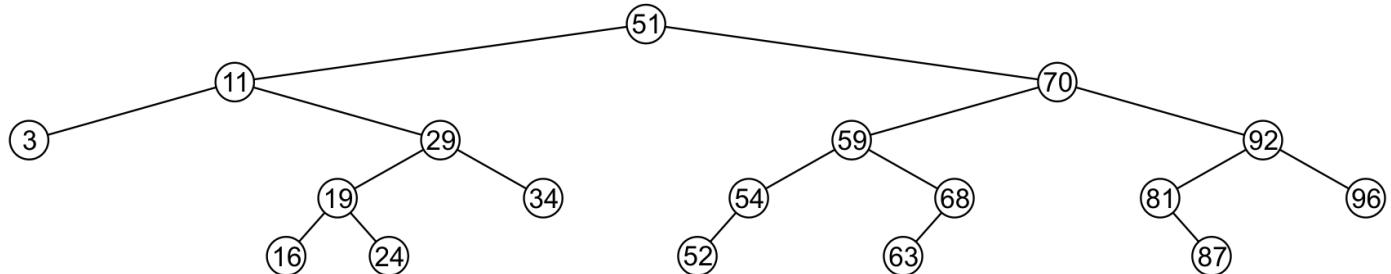


# Previous and Next Objects

More generally: find the next entry of an arbitrary object

Design a function that

- runs a single search from the root node to one of the leaves—an  $O(h)$  operation
- returns the input object if it did not find something greater than it



Ex: 0 → ; ; 25 → 29; ; 40 → 51; ; 100 → 100

# Previous and Next Objects

It returns the next object within this subtree.

Calling it from the root node return the next object in the BST.

*binary search tree*

```
template <typename Type>
```

```
Type Binary_search_node<Type>::next( Type const &obj ) const {
```

```
    if ( retrieve() == obj ) {  
        return ( right() == nullptr ) ? obj : right()->front();  
    }
```

```
    else if ( retrieve() > obj ) {  
        if( left() == nullptr )  
            return retrieve();  
        else {
```

```
            Type tmp = left()->next( obj );  
            return ( tmp == obj ) ? retrieve() : tmp;
```

↗ T .  
↓ .

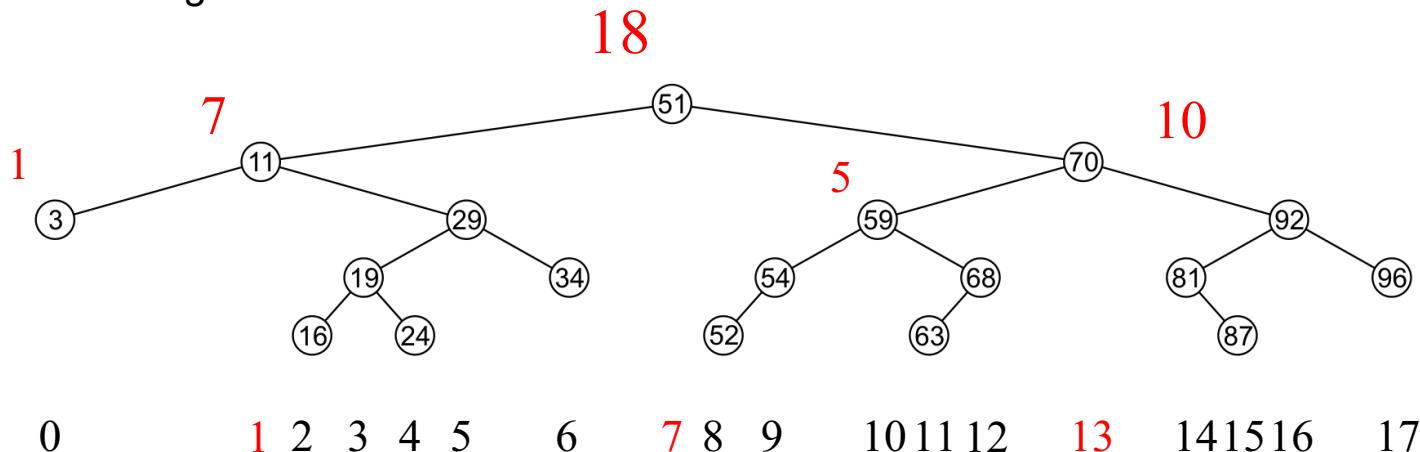
```
    } else {  
        return ( right() == nullptr ) ? obj : right()->next( obj )  
    }
```

Q :  
not in  
Sorted?

# Finding the $k^{\text{th}}$ Object

Another operation on sorted lists may be finding the  $k^{\text{th}}$  largest object

- Recall that  $k$  goes from 0 to  $n - 1$
- If the left-sub-tree has  $\ell = k$  entries, return the current node,
- If the left sub-tree has  $\ell > k$  entries, return the  $k^{\text{th}}$  entry of the left sub-tree,
- Otherwise, the left sub-tree has  $\ell < k$  entries, so return the  $(k - \ell - 1)^{\text{th}}$  entry of the right sub-tree



# Finding the $k^{\text{th}}$ Object

```
template <typename Type>
Type Binary_search_tree<Type>::at( int k ) const {
    return ( k < 0 || k >= size() ) ? nullptr : root()->at( k );
    // Need to go from 0, ..., n - 1
}
```

```
template <typename Type>
Type Binary_search_node<Type>::at( int k ) const {
    if ( left()->size() == k ) {
        return retrieve();
    } else if ( left()->size() > k ) {
        return left()->at( k );
    } else {
        return right()->at( k - left()->size() - 1 );
    }
}
```

(Here I do not check for nullptr for simplicity)

# Finding the $k^{\text{th}}$ Object

This requires that **size()** returns in (1) time

- We must have a member variable

int tree\_size;

zdt

which stores the number of descendants of this node

- This requires  $\Theta(n)$  additional memory

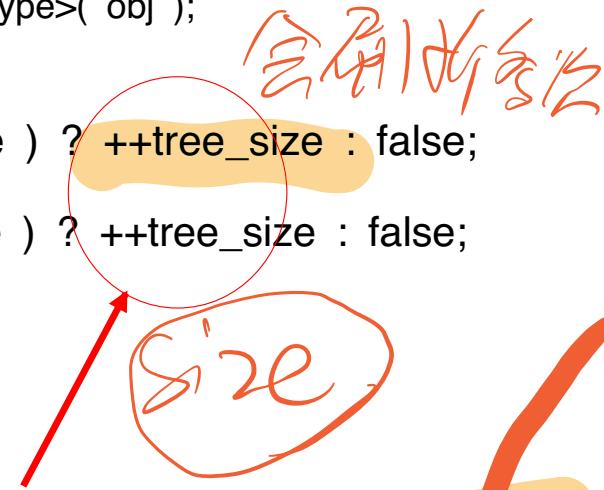
$\Theta(n)$  mem

# Finding the $k^{\text{th}}$ Object

We must now update `insert(...)` and `erase(...)` to update it

```
template <typename Type>
bool Binary_search_node<Type>::insert( Type const &obj,
                                         Binary_search_node *&ptr_to_this ) {
    if ( empty() ) {
        ptr_to_this = new Binary_search_node<Type>( obj );
        return true;
    } else if ( obj < retrieve() ) {
        return left()->insert( obj, left_tree ) ? ++tree_size : false;
    } else if ( obj > retrieve() ) {
        return right()->insert( obj, right_tree ) ? ++tree_size : false;
    } else {
        return false;
    }
}
```

Clever trick: in C and C++, any non-zero value is interpreted as true



# Summary

- Abstract Sorted Lists
  - Problems using arrays and linked lists
- Binary search tree
  - Definition
  - Implementation of:
    - Front, back, insert, erase
    - Previous smaller and next larger objects
    - Finding the  $k^{\text{th}}$  Object

# Run Time on BST

Almost all of the relevant operations on a binary search tree are  $O(h)$  *pre - sorted*.

- If the tree is *close* to a linked list, the run times is  $O(n)$ 
  - Insert 1, 2, 3, 4, 5, 6, 7, ...,  $n$  into a empty binary search tree
- The best we can do is if the tree is perfect:  $O(\ln(n))$
- Our goal will be to find tree structures where we can maintain a height of  $(\ln(n))$

## Solution

- AVL trees
- Red-black trees
- B trees, B+ trees
- Splay trees
- More...

All of which ensure that the height remains  $(\ln(n))$